

EXAMPLE REPORT · OUR OWN DATA

Does KAS v1 predict 4-hour returns?

A higher KAS score (0–100) is followed by a higher 4-hour return, across the eight coins Kvantix scores.

Report	KVX-EXAMPLE-V1
Prepared for	Public example — Kvantix's own live engine
Date	2026-09-26
Data	62,907 rows · 2026-07-06 06:42 → 2026-09-26 04:15
Horizon · cost	4h · 0.06% per round trip
Data fingerprint	SHA-256 a1947f371f5ed0a1ef0a7c0d17edfef7901d32cc08a7e8905a111815b3cdc148
Toolkit	git alea88f 2026-09-26
Prepared by	Kvantix · Signal Validation

Summary

THE CLAIM

DOES NOT HOLD

The claim does not hold. Over 62,907 scores the rank correlation between KAS and the next 4-hour return is +0.0003 — far inside the noise floor — and block-shuffled outcomes match it 99% of the time.

The average outcome does rise from the lowest to the highest third of scores, but the share of positive outcomes falls. The difference comes from a few large moves, not from a ranking. Split by market direction, high scores beat low scores only in rising markets: a long bias, not an edge. The one test v1 passes is stability — its correlation is slightly negative in both halves of the period.

We closed paid access and retracted the claims that had been built on v1. This report is published unedited, as an example of what a failed signal looks like.

THE SEVEN TESTS

1 / 7 PASSED

Indistinguishable from noise: the correlation sits inside the noise floor, and random series with the same structure do as well.

Correlation vs. noise floor — ρ +0.0003 vs. noise floor ± 0.0769	× FAIL
Monotone ranking — 52.1% → 51.2% → 50.9% positive	× FAIL
Day consistency — 24 of 75 days, $p=1.000$	× FAIL
Regime split — up +0.2460 · down -0.0713 · range -0.0825	× FAIL
Net of costs — best long-short net +0.0063% at 0.06%	× FAIL
Permutation — $p=0.991$	× FAIL
Stability — half 1 -0.0665 · half 2 -0.0256	✓ PASS

Pass or fail per test, never a single score. A claim is judged by a person, using the tests — the tests do not add up to a verdict on their own.

01 Data check

Most serious errors are found here, before any statistics: duplicated rows, outcomes that are not aligned with their timestamps, and periods or assets that are missing.

Rows	62,907 in file · 62,907 used
Period	2026-07-06 06:42 → 2026-09-26 04:15 (81.9 days)
Sampling	every 15 min (median)
Overlap	16× — each outcome window spans 16 sampling steps
Effective sample	678 independent observations — the 8 groups move together (outcome correlation 0.69), so each time step counts as 1.38 groups, not 8
Groups	8
Columns used	timestamp = timestamp · prediction = prediction · outcome = outcome · group = asset · reference = price

OK	One row per timestamp per group — no duplicated observations.
CHECK	4 outcomes lie more than 10 standard deviations from the mean. Check units (percent vs. fraction) and data errors.
OK	The outcome tracks the forward price change over 4h (rank correlation +0.99), not the past one (-0.04). Outcome and timestamp are aligned.

Group	Rows	From	To
ADA	7,863	2026-07-06	2026-09-26
AVAX	7,863	2026-07-06	2026-09-26
BNB	7,863	2026-07-06	2026-09-26
BTC	7,866	2026-07-06	2026-09-26
DOGE	7,863	2026-07-06	2026-09-26
ETH	7,863	2026-07-06	2026-09-26
SOL	7,863	2026-07-06	2026-09-26
XRP	7,863	2026-07-06	2026-09-26

HOW THE PREDICTIONS WERE PRODUCED

KAS v1 was the score Kvantix published live on kvantix.tech: one number per coin from orderbook imbalance, funding and momentum, for BTC, ETH, SOL, BNB, XRP, ADA, DOGE and AVAX.

Every published score was written to a hash-chained ledger at the timestamp shown, before its outcome existed, so the scores cannot have been revised afterwards. The outcome is the coin's return over the following 4 hours, computed later from public Binance prices by a separate process.

This export holds every ledger row with a 4-hour outcome from 6 July to 26 September 2026, taken with the same query as our internal control run.

NOTES FROM MANUAL INSPECTION

There is exactly one row per coin per 15-minute tick, and the outcome tracks the forward price change, not the past one — the two errors that most often invalidate a test are absent.

The eight coins move together: their 4-hour outcomes correlate about 0.69 on average. Eight coins at the same moment are therefore worth about 1.4 independent observations, not 8, and the effective sample and noise floor reflect that.

The four outcomes flagged as extreme fall on 22 August (ADA and XRP, +12.7% to +15.3% in four hours) and 20 September (AVAX, +14.0%). They agree with the recorded Binance prices, so they are market moves, not data errors. Rank-based tests are not distorted by them; averages are — which is why the rising average across the thirds deserves little weight.

02 The seven tests

Each test exists because it caught a signal that looked convincing. Outcomes are in % over 4h. Numbers are shown in full so every step can be checked.

1. Correlation vs. noise floor

× FAIL

Do higher predictions go with higher outcomes at all? We use rank correlation (Spearman), which is robust to outliers and to the scale of the prediction. With overlapping windows, n is not n : observations that share most of their outcome window are nearly the same observation. The noise floor is set by the effective sample size — anything inside it is indistinguishable from zero.

Spearman +0.0003 · Pearson +0.0324 · noise floor ± 0.0769 (n_{eff} 678)




Reading: The correlation is inside the noise floor: it cannot be told apart from zero.

2. Monotone ranking

× FAIL

Predictions are split into thirds. Outcomes must improve step by step from the lowest to the highest third. A signal where the middle third wins is not a signal, and a perfectly ordered but falling sequence is an inverted one.

third 1: mean +0.0363%, 52.1% positive (n 20,969) · third 2: mean +0.0593%, 51.2% positive (n 20,969) · third 3: mean +0.1315%, 50.9% positive (n 20,969)



Reading: Mixed: the average outcome rises from low to high, but the share of positive outcomes falls. The gain comes from a few large outcomes, not a broad ranking.

3. Day consistency

× FAIL

On how many days did high predictions beat low ones — and was the gap positive on average? Being right 70% of days can still lose money if the other 30% are large, so both frequency and value must hold.

24 of 75 days positive (32.0%) · binomial $p=0.9995$ · mean daily gap -0.1051% · median -0.1735%


Reading: Frequency and value do not point the same way, or neither is significant.

4. Regime split

× FAIL

Each observation is labelled by the market trend at the time of the prediction (known then, not afterwards). An edge must be positive in both rising and falling markets. Positive only while the market falls is a short bias in disguise — this test caught our own engine.

rising: gap +0.2460% (n 25,335) · falling: gap -0.0713% (n 22,996) · ranging: gap -0.0825% (n 13,808)



Reading: Positive only in rising markets — a long bias, not an edge.

5. Net of costs

× FAIL

The gap between high and low predictions, long-short, after two round trips of trading cost, at three fixed thresholds. All three are reported; none is selected afterwards. It must beat the cost by a margin of half a round trip. Long-only results do not count: they also measure the market's direction.

cost 0.06% per round trip · required margin 0.0300% · top/bottom 40%: gap +0.0688, long-short net -0.0512 · top/bottom 30%: gap +0.1131, long-short net -0.0069 · top/bottom 20%: gap +0.1263, long-short net +0.0063

Threshold	Coverage	High	Low	Gap	Long-only net	Long-short net
top/bottom 40%	80%	+0.1116	+0.0428	+0.0688	+0.0516	-0.0512
top/bottom 30%	60%	+0.1467	+0.0337	+0.1131	+0.0867	-0.0069
top/bottom 20%	40%	+0.1727	+0.0463	+0.1263	+0.1127	+0.0063


Reading: No threshold clears the cost with margin.

6. Permutation

× FAIL

Outcomes are shuffled in blocks — preserving their structure and autocorrelation, breaking only the link to the predictions — many times. If random series produce effects as large as yours, the effect is not evidence of skill.

observed p +0.0003 · random median -0.0003 · random 95th percentile +0.0395 · p=0.9910 · 1,000 permutations, block size 128




Reading: Random series with the same structure do as well — the effect is not evidence of skill.

7. Stability

✓ PASS

The period is split in two halves. A real relationship has the same sign in both. One that only exists in one half is a period effect.

half 1 (2026-07-06 → 2026-08-16, n 31,453): ρ -0.0665 · half 2 (2026-08-16 → 2026-09-26, n 31,454): ρ -0.0256



Reading: Same sign in both halves.

03 Further analysis

Cost sensitivity

Long-short net result per trade at different round-trip costs, for each of the three fixed thresholds. ✓ marks results that clear the toolkit's margin (half a round trip). The client's stated cost is highlighted.

Round-trip cost	top/bottom 40%	top/bottom 30%	top/bottom 20%
0.00%	+0.0688% ✓	+0.1131% ✓	+0.1263% ✓
0.02%	+0.0288% ✓	+0.0731% ✓	+0.0863% ✓
0.04%	-0.0112%	+0.0331% ✓	+0.0463% ✓
0.06%	-0.0512%	-0.0069%	+0.0063%
0.08%	-0.0912%	-0.0469%	-0.0337%
0.10%	-0.1312%	-0.0869%	-0.0737%
0.15%	-0.2312%	-0.1869%	-0.1737%
0.20%	-0.3312%	-0.2869%	-0.2737%

The most the signal can bear at any of the three thresholds, with the margin, is **0.0505%** per round trip. Break-even without margin: 0.0344% (top/bottom 40%), 0.0566% (top/bottom 30%), 0.0631% (top/bottom 20%).

By group — none excluded

4 of 8 groups have a positive rank correlation. The effect is spread out. Selecting the groups that look best after the fact is exactly how a period effect gets mistaken for an edge.

Group	Rows	Spearman
ADA	7,863	+0.0184
AVAX	7,863	-0.0153
BNB	7,863	+0.0133
BTC	7,866	-0.0145
DOGE	7,863	+0.0095
ETH	7,863	-0.0274
SOL	7,863	-0.0065
XRP	7,863	+0.0133

Outcome alignment

From the price column we computed the change over 4h forward and backward from every timestamp (62,628 rows). Rank correlation of the outcome with the forward change: +0.991; with the past change: -0.038. The prediction's correlation with the past change is +0.353 — how much of it is momentum or reversal.

ADDITIONAL ANALYSIS

What KAS v1 was actually measuring. Its score correlates +0.35 (rank) with the coin's price change over the previous four hours, and +0.0003 with the next four. In practice v1 was a reading of recent momentum — which told us where a coin had been, not where it was going.

04 Verdict on the claim

A higher KAS score (0–100) is followed by a higher 4-hour return, across the eight coins Kvantix scores.

DOES NOT HOLD

REASONING

The claim does not hold. Over 62,907 scores the rank correlation between KAS and the next 4-hour return is +0.0003 — far inside the noise floor — and block-shuffled outcomes match it 99% of the time.

The average outcome does rise from the lowest to the highest third of scores, but the share of positive outcomes falls. The difference comes from a few large moves, not from a ranking. Split by market direction, high scores beat low scores only in rising markets: a long bias, not an edge. The one test v1 passes is stability — its correlation is slightly negative in both halves of the period.

We closed paid access and retracted the claims that had been built on v1. This report is published unedited, as an example of what a failed signal looks like.

05 What it would take

What the data would have to show for the claim to be made honestly. Derived from the tests above; figures assume the effect stays as measured.

- 1 An effect to measure.** The correlation ($\rho = +0.0003$) is essentially zero. At that size it would take more than a century more of data to separate it from noise. More data of the same kind will not rescue the claim.
- 2 A gap that survives costs.** At 0.06% per round trip, the gap between high and low predictions must exceed 0.150% (two round trips plus a margin of half a round trip). The largest gap at any of the three fixed thresholds is +0.1263%. Equivalently, the signal can bear at most 0.0505% per round trip with that margin.
- 3 An edge in both market directions.** The high-minus-low gap must be positive in both rising and falling markets. Positive in one direction only is a directional bias.
- 4 Beat random series with the same structure.** Block-shuffled outcomes produce effects as large as the observed one. Part of what looks like skill is the structure of the data.
- 5 Consistency that also pays.** High predictions must beat low ones on a significant share of days AND on average. Frequent small wins with rare large losses do not count.
- 6 A pre-registered forward test.** Whatever the result here, the only convincing evidence is a claim that is written down and locked before the data exists, then measured on that data alone. Kvantix Forward Test does exactly that.

06 Method and limitations

The toolkit knows only three things about each row: when the prediction was made, what it was, and what happened over the horizon. It knows nothing about trading, and it does not search for parameters — every threshold is fixed in advance, and all of them are reported.

Effective sample. When outcomes are measured over a horizon longer than the sampling interval, neighbouring observations share most of their window, so the row count is divided by the overlap. When several groups (assets, segments) are observed at the same time and their outcomes move together, they are not independent either: each time step counts as $G / (1 + (G - 1) \cdot \rho)$ groups, where ρ is the measured average correlation between the groups' outcomes. All significance statements use the resulting effective sample.

Permutation. Outcomes are rotated and shuffled in blocks that span one full outcome window in time (across all groups), which preserves their autocorrelation and cross-group structure and breaks only their link to the predictions. A naive shuffle would understate the noise.

What the tests cannot see. Lookahead that is baked into the data (a prediction computed with information from after its timestamp) looks like skill to any statistical test. The data check looks for its symptoms, but the provenance of the predictions has to be established with the client. Likewise, a result on past data is evidence about the past; only a pre-registered forward test says something about the future.

Units. Outcomes are used as delivered. Net-of-costs figures assume outcomes and costs are in the same unit (percent).

Permutations: 1,000 (block size 128). Noise floor = $2/\sqrt{n_{\text{eff}}-1} = 0.0769$. Regime lookback: 96 observations per group, trend threshold $\pm 0.5\%$.

This report assesses statistics, not whether anyone should trade. It is not investment advice. If quoted, it must be quoted in full — not only the tests that passed. Client data is deleted 30 days after delivery. Kvantix · CVR 46296036 · Denmark.