

EXAMPLE REPORT · OUR OWN DATA

Is KAS v2.1 tradeable?

KAS v2.1 ranks the next 4-hour returns of the eight coins well enough to trade at 0.06% per round trip.

Report	KVX-EXAMPLE-V2
Prepared for	Public example — Kvantix's own shadow engine
Date	2026-09-26
Data	47,944 rows · 2026-07-14 05:45 → 2026-09-14 15:45
Horizon · cost	4h · 0.06% per round trip
Data fingerprint	SHA-256 c79732f702943dd28b6d1f5509c35a15706997bedc2c15217cd4f618cdd14d68
Toolkit	git alea88f 2026-09-26
Prepared by	Kvantix · Signal Validation

Summary

THE CLAIM DOES NOT HOLD The claim does not hold, although v2.1 got further than v1. High scores beat low scores on 45 of 63 days ($p = 0.0004$), and the share of positive outcomes rises steadily from the lowest to the highest third. That is a consistent pattern. It is not an edge. Over the whole period the rank correlation (+0.025) is inside the noise floor, block-shuffled outcomes reach it about one time in ten, it is positive in the first half and slightly negative in the second, and positive only in falling markets. Most decisively, the gap between the highest- and lowest-scored coins is negative at all three thresholds — before any trading cost. A signal that is right on most days but loses on the days that matter cannot be traded. We stopped the shadow on 14 September. Nothing built on it was ever offered for sale.	THE SEVEN TESTS 2 / 7 PASSED Indistinguishable from noise: the correlation sits inside the noise floor, and random series with the same structure do as well.
---	--

Monotone ranking — 49.7% → 50.5% → 53.5% positive	✓ PASS
Day consistency — 45 of 63 days, p=0.000	✓ PASS
Regime split — up -0.0718 · down +0.0361 · range -0.0163	× FAIL
Net of costs — best long-short net -0.1432% at 0.06%	× FAIL
Permutation — p=0.098	× FAIL
Stability — half 1 +0.0600 · half 2 -0.0111	× FAIL

Pass or fail per test, never a single score. A claim is judged by a person, using the tests — the tests do not add up to a verdict on their own.

01 Data check

Most serious errors are found here, before any statistics: duplicated rows, outcomes that are not aligned with their timestamps, and periods or assets that are missing.

Rows	47,944 in file · 47,944 used
Period	2026-07-14 05:45 → 2026-09-14 15:45 (62.4 days)
Sampling	every 15 min (median)
Overlap	16× — each outcome window spans 16 sampling steps
Effective sample	509 independent observations — the 8 groups move together (outcome correlation 0.70), so each time step counts as 1.36 groups, not 8
Groups	8
Columns used	timestamp = timestamp · prediction = prediction · outcome = outcome · group = asset · reference = price

OK	One row per timestamp per group — no duplicated observations.
CHECK	5 outcomes lie more than 10 standard deviations from the mean. Check units (percent vs. fraction) and data errors.
OK	The outcome tracks the forward price change over 4h (rank correlation +0.99), not the past one (-0.07). Outcome and timestamp are aligned.

Group	Rows	From	To
ADA	5,993	2026-07-14	2026-09-14
AVAX	5,993	2026-07-14	2026-09-14
BNB	5,993	2026-07-14	2026-09-14
BTC	5,993	2026-07-14	2026-09-14
DOGE	5,993	2026-07-14	2026-09-14
ETH	5,993	2026-07-14	2026-09-14
SOL	5,993	2026-07-14	2026-09-14
XRP	5,993	2026-07-14	2026-09-14

HOW THE PREDICTIONS WERE PRODUCED

KAS v2.1 was a rebuilt engine that ran in shadow from 14 July to 14 September 2026: computed and logged live, but never published, traded or shown to anyone — including us — until this evaluation.

The shadow logged a score every 60 seconds. For each 15-minute tick in the public ledger we take the shadow score logged closest to it (within 7.5 minutes) and the same 4-hour outcome that v1 was measured against, computed from public Binance prices.

NOTES FROM MANUAL INSPECTION

Our first extraction of this data matched every 60-second shadow row to every 15-minute tick. The same underlying data came out as 145,800 rows instead of about 54,000, the noise floor fell from 0.034 to 0.021, and the verdict changed. Duplicated rows are the reason the data check comes first. This export has exactly one row per coin per tick.

As for v1, the eight coins' outcomes move together (average correlation about 0.70), so each tick is worth about 1.4 independent observations.

The five outcomes flagged as extreme are the same market event as in the v1 report: ADA and XRP on 22 August (between -11.9% and +15.3% in four hours). They agree with the recorded Binance prices. Rank-based tests are not distorted by them; averages are.

02 The seven tests


Each test exists because it caught a signal that looked convincing. Outcomes are in % over 4h. Numbers are shown in full so every step can be checked.

1. Correlation vs. noise floor

× FAIL

Do higher predictions go with higher outcomes at all? We use rank correlation (Spearman), which is robust to outliers and to the scale of the prediction. With overlapping windows, n is not n : observations that share most of their outcome window are nearly the same observation. The noise floor is set by the effective sample size — anything inside it is indistinguishable from zero.

Spearman +0.0247 · Pearson -0.0102 · noise floor ± 0.0887 (n_{eff} 509)




Reading: The correlation is inside the noise floor: it cannot be told apart from zero.

2. Monotone ranking

✓ PASS

Predictions are split into thirds. Outcomes must improve step by step from the lowest to the highest third. A signal where the middle third wins is not a signal, and a perfectly ordered but falling sequence is an inverted one.

third 1: mean +0.0964%, 49.7% positive (n 15,981) · third 2: mean +0.0548%, 50.5% positive (n 15,981) · third 3: mean +0.0682%, 53.5% positive (n 15,982)



Reading: The share of positive outcomes rises step by step; the average outcome does not.

3. Day consistency

✓ PASS

On how many days did high predictions beat low ones — and was the gap positive on average? Being right 70% of days can still lose money if the other 30% are large, so both frequency and value must hold.

45 of 63 days positive (71.4%) · binomial $p=0.0004$ · mean daily gap +0.1231% · median +0.1132%


Reading: High beats low on a significant share of days, and on average.

4. Regime split

× FAIL

Each observation is labelled by the market trend at the time of the prediction (known then, not afterwards). An edge must be positive in both rising and falling markets. Positive only while the market falls is a short bias in disguise — this test caught our own engine.

rising: gap -0.0718% (n 18,144) · falling: gap $+0.0361\%$ (n 17,440) · ranging: gap -0.0163% (n 11,592)



Reading: Positive only in falling markets — a short bias, not an edge.

5. Net of costs

× FAIL

The gap between high and low predictions, long-short, after two round trips of trading cost, at three fixed thresholds. All three are reported; none is selected afterwards. It must beat the cost by a margin of half a round trip. Long-only results do not count: they also measure the market's direction.

cost 0.06% per round trip · required margin 0.0300% · top/bottom 40%: gap -0.0232 , long-short net -0.1432 · top/bottom 30%: gap -0.0335 , long-short net -0.1535 · top/bottom 20%: gap -0.0387 , long-short net -0.1587

Threshold	Coverage	High	Low	Gap	Long-only net	Long-short net
top/bottom 40%	85%	+0.0656	+0.0888	-0.0232	+0.0056	-0.1432
top/bottom 30%	63%	+0.0639	+0.0973	-0.0335	+0.0039	-0.1535
top/bottom 20%	44%	+0.0589	+0.0976	-0.0387	-0.0011	-0.1587


Reading: No threshold clears the cost with margin.

6. Permutation

× FAIL

Outcomes are shuffled in blocks — preserving their structure and autocorrelation, breaking only the link to the predictions — many times. If random series produce effects as large as yours, the effect is not evidence of skill.

observed p $+0.0247$ · random median $+0.0002$ · random 95th percentile $+0.0259$ · $p=0.0979$ · 1,000 permutations, block size 128




Reading: Random series with the same structure do as well — the effect is not evidence of skill.

7. Stability

× FAIL

The period is split in two halves. A real relationship has the same sign in both. One that only exists in one half is a period effect.

half 1 (2026-07-14 → 2026-08-14, n 23,972): ρ +0.0600 · half 2 (2026-08-14 → 2026-09-14, n 23,972): ρ -0.0111



Reading: The sign changes between the halves.

03 Further analysis

Cost sensitivity

Long-short net result per trade at different round-trip costs, for each of the three fixed thresholds. ✓ marks results that clear the toolkit's margin (half a round trip). The client's stated cost is highlighted.

Round-trip cost	top/bottom 40%	top/bottom 30%	top/bottom 20%
0.00%	-0.0232%	-0.0335%	-0.0387%
0.02%	-0.0632%	-0.0735%	-0.0787%
0.04%	-0.1032%	-0.1135%	-0.1187%
0.06%	-0.1432%	-0.1535%	-0.1587%
0.08%	-0.1832%	-0.1935%	-0.1987%
0.10%	-0.2232%	-0.2335%	-0.2387%
0.15%	-0.3232%	-0.3335%	-0.3387%
0.20%	-0.4232%	-0.4335%	-0.4387%

The gap between high and low predictions is zero or negative at every threshold, so no level of cost — not even zero — makes it profitable.

By group — none excluded

6 of 8 groups have a positive rank correlation. The effect is broad.

Group	Rows	Spearman
ADA	5,993	+0.0499
AVAX	5,993	+0.0168
BNB	5,993	-0.0309
BTC	5,993	+0.0304
DOGE	5,993	+0.0484
ETH	5,993	+0.0705
SOL	5,993	+0.0301
XRP	5,993	-0.0073

Outcome alignment

From the price column we computed the change over 4h forward and backward from every timestamp (47,688 rows). Rank correlation of the outcome with the forward change: +0.991; with the past change: -0.073. The prediction's correlation with the past change is -0.373 — how much of it is momentum or reversal.

ADDITIONAL ANALYSIS

What changed from v1. v2.1's score correlates -0.37 (rank) with the previous four-hour price change, where v1 correlated +0.35: the rebuilt engine leaned against recent moves instead of following them. That produced the day-to-day consistency — and still no usable gap at the extremes, where a trade would actually be placed.

04 Verdict on the claim

KAS v2.1 ranks the next 4-hour returns of the eight coins well enough to trade at 0.06% per round trip.

DOES NOT HOLD

REASONING

The claim does not hold, although v2.1 got further than v1. High scores beat low scores on 45 of 63 days ($p = 0.0004$), and the share of positive outcomes rises steadily from the lowest to the highest third. That is a consistent pattern.

It is not an edge. Over the whole period the rank correlation (+0.025) is inside the noise floor, block-shuffled outcomes reach it about one time in ten, it is positive in the first half and slightly negative in the second, and positive only in falling markets. Most decisively, the gap between the highest- and lowest-scored coins is negative at all three thresholds — before any trading cost. A signal that is right on most days but loses on the days that matter cannot be traded.

We stopped the shadow on 14 September. Nothing built on it was ever offered for sale.

05 What it would take

What the data would have to show for the claim to be made honestly. Derived from the tests above; figures assume the effect stays as measured.

- 1 **Enough independent observations.** At the observed effect size ($p = +0.0247$), the correlation would clear the noise floor at an effective sample of about 6,557 (now 509). At the current pace that is roughly 2 years more of data — and only if the effect is real and stable, which more data may equally show it is not.
- 2 **A gap that survives costs.** At 0.06% per round trip, the gap between high and low predictions must exceed 0.150% (two round trips plus a margin of half a round trip). The largest gap at any of the three fixed thresholds is -0.0232%. The gap is negative at every threshold: the highest predictions do worse than the lowest before any cost is paid.
- 3 **An edge in both market directions.** The high-minus-low gap must be positive in both rising and falling markets. Positive in one direction only is a directional bias.
- 4 **The same sign in both halves.** The relationship reverses between the first and second half of the period. A claim needs it to hold across periods.
- 5 **Beat random series with the same structure.** Block-shuffled outcomes produce effects as large as the observed one. Part of what looks like skill is the structure of the data.
- 6 **A pre-registered forward test.** Whatever the result here, the only convincing evidence is a claim that is written down and locked before the data exists, then measured on that data alone. Kvantix Forward Test does exactly that.

06 Method and limitations

The toolkit knows only three things about each row: when the prediction was made, what it was, and what happened over the horizon. It knows nothing about trading, and it does not search for parameters — every threshold is fixed in advance, and all of them are reported.

Effective sample. When outcomes are measured over a horizon longer than the sampling interval, neighbouring observations share most of their window, so the row count is divided by the overlap. When several groups (assets, segments) are observed at the same time and their outcomes move together, they are not independent either: each time step counts as $G / (1 + (G - 1) \cdot \rho)$ groups, where ρ is the measured average correlation between the groups' outcomes. All significance statements use the resulting effective sample.

Permutation. Outcomes are rotated and shuffled in blocks that span one full outcome window in time (across all groups), which preserves their autocorrelation and cross-group structure and breaks only their link to the predictions. A naive shuffle would understate the noise.

What the tests cannot see. Lookahead that is baked into the data (a prediction computed with information from after its timestamp) looks like skill to any statistical test. The data check looks for its symptoms, but the provenance of the predictions has to be established with the client. Likewise, a result on past data is evidence about the past; only a pre-registered forward test says something about the future.

Units. Outcomes are used as delivered. Net-of-costs figures assume outcomes and costs are in the same unit (percent).

Permutations: 1,000 (block size 128). Noise floor = $2/\sqrt{n_{\text{eff}}-1}$ = 0.0887. Regime lookback: 96 observations per group, trend threshold $\pm 0.5\%$.

KVX-EXAMPLE-V2 · SHA-256 of input data c79732f702943dd28b6d1f5509c35a15706997bedc2c15217cd4f618cdd14d68 · git a1ea88f 2026-09-26

This report assesses statistics, not whether anyone should trade. It is not investment advice. If quoted, it must be quoted in full — not only the tests that passed. Client data is deleted 30 days after delivery. Kvantix · CVR 46296036 · Denmark.